

# Diagonalizable Matrices

- A matrix  $R$  is *diagonalizable* when it can be written as a product shown below
- $P$  is a non-singular matrix
- Called a *similarity transformation*
  - How an operator changes under a change of basis

$$R = P^{-1}AP \tag{1}$$

# Define Polynomials of Matrices

- Polynomial Algebra of matrices:
  - 1 corresponds to identity
  - scalar multiplication
  - powers are products of matrix with itself or iterated compositions
- Cayley-Hamilton theorem: any matrix satisfies its own characteristic polynomial

$$f(A) = a_n A^n + \cdots + a_1 A + a_0 I \quad (2)$$

# Characteristic Polynomial

- The Characteristic Polynomial of a matrix is computed by means of a special determinant
- The roots of characteristic polynomial are the *eigenvalues*
- Eigenvectors are vectors satisfying the expression

$$\Delta(\lambda) = |(\lambda I - A)| \quad (3)$$

$$Rv = \lambda v \quad (4)$$

$$0 = (\lambda I - R)v \quad (5)$$

# Diagonal Factorization, Spectral Decomposition


- Reversing the process we can write a matrix as a similarity transformation of a diagonal matrix
- If  $A$  is real-symmetric then the roots of its characteristic polynomial are guaranteed to be real.
- Eigenvectors belonging to distinct eigenvalues are orthogonal to one another  $u \cdot v = 0$

$$A = PDP^{-1} \tag{6}$$

# Principal Component Analysis in Two Dimensions

## Box 1

A (supervised) ML workflow will minimally contain: *data, feature selection process, cross-validation, an ML algorithm, and a means to evaluate performance*



# Model Evaluation

- Would like to assess, quantitatively, how well our algorithm can make predictions
- *Resubstitution Loss*: Performance of the classifier on the training data set
- Receiver Operating Characteristic (ROC): Visualize how the sensitivity and specificity of a test changes as some threshold parameter is varied

# Cross-Validation

- Separation of Data into Training and Test sets
- k-Fold: The data-set is split into  $k$  separate partitions, each partition is called a *fold*
- We perform our evaluation  $k$  times, withholding a different fold from the training set in each case
- The performance of the result is reported as the average over all of the separate folds
- *leave-one-out* CV is a limiting case where each fold is a singleton


# Ensemble Methods

- Goal: Combine a number of weak-predictors into a single strong-predictor
- *Bagging* (Bootstrap Aggregating): Train the same machine learning algorithm on a set of distinct bootstrap samples and combine the results by majority voting
- Make a predictor more robust by getting a variety of different views of a dataset



# Basic Two-Outcome Tests

- Purpose of machine learning is to create some function in feature space that allows us to classify data points as belonging to some known classes
- Confusion matrix is used to characterize the different types of error that can occur



# Performance Metrics

- In any binary classification task two types of error can occur, false positives and false negatives
- In general we need to keep track of both of these errors to understand how well our classifier is performing
- The *Receiver Operating Characteristic* (ROC) keeps track of both of these error rates as we vary the threshold

$$\text{threshold} > w^T x \quad (7)$$

## Different Types of Error, Classifier Performance

$$\text{Confusion Matrix} = \begin{bmatrix} TP & FN \\ FP & TN \end{bmatrix} \quad (8)$$


- Accuracy:  $(TP + TN) / \text{TOTAL}$
- Sensitivity:  $TP / (TP + FN)$
- Specificity:  $TN / (FP + TN)$

# Making Distinction Between Covariance and Confusion

- The *covariance matrix* describes how two quantities vary in relation to one another
- The diagonal terms are simply the variance of that particular variable and the off diagonal correspond to each pair
- The *confusion matrix* enters when we have a classification problem, it is a metric of how well some classification procedure performed compared to known values

# An example

- Typical classification problem
- In this case, data consists of ordered pairs such that  $(\text{class label} \in (0, 1), \text{variate})$
- Various calculations we can do with this data



## We can look at the covariance and correlation


- Measure the natural tendency of these two numbers to vary together
- In this case, a p-value can be recovered from the correlation and is of the order  $10^{-24}$

$$R_{cov} = \begin{bmatrix} 1.1477 & -0.2549 \\ -0.2549 & 0.2506 \end{bmatrix} \quad (9)$$

$$R_{corr} = \begin{bmatrix} 1.0000 & -0.4752 \\ -0.4752 & 1.0000 \end{bmatrix} \quad (10)$$


## We can also examine the histogram

- With enough data points, the histogram becomes smooth and we can see how each individual class contributes to the total distribution
- The means are clearly different, but a simple threshold classifier won't perform well since there is a lot of overlap





# Receiver Operating Characteristic

- The Receiver Operating Characteristic (ROC) summarizes all possible threshold tests and presents in a space representing sensitivity and specificity
- The *area under the curve*, in this case  $\approx 0.77$  is another measure of the relationship between the variate and the class label
- What would a perfect or near-perfect test look like?






# More Favorable Data



# More Favorable Data ROC



# Steps for Performing LDA

- Linear Discriminant Analysis (Binary)
  - Compute the within class means  $\mu_1$  and  $\mu_2$
  - Compute sample covariances in each class  $S_1$  and  $S_2$
  - Compute the within class scatter matrix  $S_W = S_1 + S_2$
  - Find the discriminant projection  $w$  according to the equation (5) below
  - Choose a value for the threshold

$$\hat{w} = S_W^{-1}(\mu_1 - \mu_2) \quad (11)$$

$$\text{threshold} > \hat{w}^T x \quad (12)$$

# Calculations

- Suppose we want to find the covariances for the datasets below
- Use a library or do calculation manually

$$\mathbf{A} = \begin{bmatrix} 0 & 0 \\ 0 & 0 \\ 0 & 0 \\ 0 & 0 \\ 0 & 0 \\ 0 & 0 \\ 0 & 0 \end{bmatrix}, \mathbf{B} = \begin{bmatrix} 0 & 0 \\ 0 & 0 \\ 0 & 0 \\ 0 & 1 \\ 0 & 0 \\ 0 & 0 \\ 0 & 0 \end{bmatrix}, \mathbf{C} = \begin{bmatrix} 0 & 0 \\ 0 & 0 \\ 2 & 0 \\ 0 & 0 \\ 0 & 0 \\ 0 & 0 \\ 0 & 0 \end{bmatrix}, \mathbf{D} = \begin{bmatrix} 0 & 1 \\ 1 & 0 \\ 0 & 1 \\ 1 & 0 \\ 0 & 1 \\ 1 & 0 \\ 1 & 0 \end{bmatrix}, \quad (13)$$

# General Solution

- Suppose we want to find the covariances for the datasets below
- Use a library or do calculation manually

$$R_{ij} = \sigma_{ij}^2 = E[(X_i - \mu_{x_i})(X_j - \mu_{x_j})] \quad (14)$$

$$\frac{1}{5} \left( X - \begin{bmatrix} 1 \\ 1 \\ 1 \\ 1 \\ 1 \end{bmatrix} [\mu_1, \mu_2] \right)^T \left( X - \begin{bmatrix} 1 \\ 1 \\ 1 \\ 1 \\ 1 \end{bmatrix} [\mu_1, \mu_2] \right) \quad (15)$$

## Cases A and B

Case A:

$$\frac{1}{5} \left( \begin{pmatrix} \begin{bmatrix} 0 & 0 \\ 0 & 0 \\ 0 & 0 \\ 0 & 0 \\ 0 & 0 \end{bmatrix} - \begin{bmatrix} 1 \\ 1 \\ 1 \\ 1 \\ 1 \end{bmatrix} [0, 0] \end{pmatrix}^T \begin{pmatrix} \begin{bmatrix} 0 & 0 \\ 0 & 0 \\ 0 & 0 \\ 0 & 0 \\ 0 & 0 \end{bmatrix} - \begin{bmatrix} 1 \\ 1 \\ 1 \\ 1 \\ 1 \end{bmatrix} [0, 0] \end{pmatrix} = R_{A,ij} = \begin{bmatrix} 0 & 0 \\ 0 & 0 \end{bmatrix} \quad (16)$$

Case B:

$$\frac{1}{5} \left( \begin{pmatrix} \begin{bmatrix} 0 & 0 \\ 0 & 0 \\ 0 & 0 \\ 0 & 1 \\ 0 & 0 \end{bmatrix} - \begin{bmatrix} 1 \\ 1 \\ 1 \\ 1 \\ 1 \end{bmatrix} [0, 1/6] \end{pmatrix}^T \begin{pmatrix} \begin{bmatrix} 0 & 0 \\ 0 & 0 \\ 0 & 0 \\ 0 & 1 \\ 0 & 0 \end{bmatrix} - \begin{bmatrix} 1 \\ 1 \\ 1 \\ 1 \\ 1 \end{bmatrix} [0, 1/6] \end{pmatrix} = R_{B,ij} = \begin{bmatrix} 0 & 0 \\ 0 & 1/6 \end{bmatrix} \quad (17)$$

## Cases C and D

- Sample covariances for cases C and D

Case C:

$$R_{C,ij} = \begin{bmatrix} 2/3 & 0 \\ 0 & 0 \end{bmatrix} \quad (18)$$

Case D:

$$R_{D,ij} = \begin{bmatrix} 3/10 & -3/10 \\ -3/10 & 3/10 \end{bmatrix} \quad (19)$$

# Task

- Write a python script to verify these results
- Calculate covariance matrices by hand and also with `numpy.cov`



# Matrix Calculus

- Gradient descent for machine learning is most often carried out over *vector* arguments
- In this case we just differentiate component by component and sometimes this allows us to collect terms in a compact notation
- See some familiar derivative identities compared to single-variable calculus

$$\frac{\partial \mathbf{x}^T A \mathbf{x}}{\partial \mathbf{x}} = \mathbf{x}^T (A + A^T) \quad (20)$$

# Matrix Calculus: Differentiating a scalar by a vector

- For some scalar quantity  $u(\mathbf{x})$  that depends on a vector  $\mathbf{x}$  the derivative is another vector of the partials with respect to each component

$$\frac{\partial u}{\partial \mathbf{x}} = \left[ \frac{\partial u}{\partial x_1}, \frac{\partial u}{\partial x_2}, \frac{\partial u}{\partial x_3}, \dots, \frac{\partial u}{\partial x_n} \right] \quad (21)$$

- What is the derivative of a vector with respect to another vector?

## Matrix Calculus: chain and product rules

- Same basic rules of differentiation apply when we differentiate with respect to matrices or vectors

$$\frac{\partial \mathbf{u} \cdot \mathbf{v}}{\partial \mathbf{x}} = ? \quad (22)$$

- Derivative of a dot product of two vectors  $\mathbf{u}$  and  $\mathbf{v}$  that may depend on a third  $\mathbf{x}$ ? What rule might we want to apply here?

## Matrix notation for large

- A practical machine learning problem may involve vectors and matrices of large dimension
- Matrix notation allows us a convenient method of keeping track of all of the independent parameters and express the update in a convenient fashion
- For example to do least squares for a simple matrix problem we find  $\|\mathbf{Ax} - \mathbf{b}\|^2$

$$\nabla \text{ w.r.t } \mathbf{x} = 2\mathbf{A}^T(\mathbf{Ax} - \mathbf{b}) \quad (23)$$